

# MA4106: Statistics II

Satvik Saha, 19MS154

November 13, 2022

## 1 Model selection

We return to the Boston Housing dataset, and examine the variation of  $\text{LOGMEDV}$  (the logarithm of the median value of homes in \$1000's) with 13 related variables. If we are willing to perform linear regression using only a subset of our independent variables, there are  $2^{13} = 8192$  models to consider and compare. Here, we use the Akaike Information Criterion (AIC) to compare two models, calculated as

$$\text{AIC} = n \log \left( \frac{\text{RSS}}{n} \right) + 2K,$$

where the given model has  $K$  parameters, fitted on a dataset of size  $n$ , and gives a residual sum of squares error RSS. Better models have lower AIC.

*Remark: In this particular case, our dataset is fairly small, allowing us to brute force all 8192 models relatively quickly. A summary of the best 5 models (by AIC) is given below. Here,  $\Delta\text{AIC}$  denotes the difference between the AIC and the smallest AIC in the list. Note that the model with all 13 parameters ranks fifth.*

Table 1: Best models (brute force)

$K$	$R^2$	$R_{adj}^2$	$\Delta\text{AIC}$	Features
11	0.7891	0.7844	0.0000	CRIM, ZN, CHAS, NOX, RM, DIS, RAD, TAX, PTRATIO, B, LSTAT
12	0.7896	0.7845	0.9676	CRIM, ZN, INDUS, CHAS, NOX, RM, DIS, RAD, TAX, PTRATIO, B, LSTAT
12	0.7892	0.7841	1.8363	CRIM, ZN, CHAS, NOX, RM, AGE, DIS, RAD, TAX, PTRATIO, B, LSTAT
10	0.7874	0.7831	2.1086	CRIM, CHAS, NOX, RM, DIS, RAD, TAX, PTRATIO, B, LSTAT
13	0.7896	0.7841	2.8044	CRIM, ZN, INDUS, CHAS, NOX, RM, AGE, DIS, RAD, TAX, PTRATIO, B, LSTAT

We can also adopt the following stepwise procedures.

- Forward selection:** Start with a model without any independent variables, and create  $K$  models each with a single parameter  $\beta_i$ . For each, test the hypothesis  $\beta_i \neq 0$  against the null hypothesis  $\beta_i = 0$  using an  $F$ -test, and choose the model with the best (highest)  $F$ -score (indicating that the corresponding  $\beta_i$  is significant). Repeat the above, adding parameters one by one until the  $F$ -scores are not significantly high, i.e. all  $F$ -scores are below some threshold  $F_{in}$ .
- Backward elimination:** Start with a model with all  $K$  independent variables, and create  $K$  models each removing a single parameter  $\beta_i$ . For each, test the hypothesis  $\beta_i \neq 0$  against the null hypothesis  $\beta_i = 0$  using an  $F$ -test, and choose the model with the worst (lowest)  $F$ -score (indicating that the corresponding  $\beta_i$  is not significant). Repeat the above, removing parameters one by one until the  $F$ -scores become significantly high, i.e. all  $F$ -scores are above some threshold  $F_{out}$ .
- Forward-backward:** Start with an empty model, then perform alternating forward selection and backward elimination steps, adding and removing variables based on thresholds  $F_{in}$  and  $F_{out}$ .

Below are the results of these three selection procedures, using thresholds  $F_{out} = 1.0, F_{in} = 2.0$ .

Table 2: Best models (stepwise selection)

Procedure	$K$	$R^2_{adj}$	$\Delta AIC$	Features
Forward	11	0.7844	0.0000	LSTAT, PTRATIO, CRIM, RM, DIS, NOX, B, RAD, TAX, CHAS, ZN
Forward-Backward	11	0.7844	0.0000	LSTAT, PTRATIO, CRIM, RM, DIS, NOX, B, RAD, TAX, CHAS, ZN
Backward	12	0.7845	0.9676	DIS, PTRATIO, CRIM, CHAS, RAD, B, TAX, RM, LSTAT, ZN, INDUS, NOX
OLS	13	0.7841	2.8044	CRIM, ZN, INDUS, CHAS, NOX, RM, AGE, DIS, RAD, TAX, PTRATIO, B, LSTAT

We see that both the forward and the forward-backward procedures select the ‘best’ model with 11 variables. Note that the features for these two rows in the table are presented in order of addition. The backward elimination procedure produces the next best model with 12 variables; the variable `INDUS` is not eliminated. *This model has a slightly better  $R^2_{adj}$  than the remaining ones.*

Here is the output for the backward elimination process. Each line represents a model with the indicated variable removed. Here, `D_AIC` represents the difference between the AIC of the new model with the existing model; a negative value indicates an improvement.

```

Backward Elimination
Current model uses features ['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS',
                             'RAD', 'TAX', 'PTRATIO', 'B', 'LSTAT']
-AGE      => F = 0.1587, p = 0.6906, adjR2 = 0.7845, D_AIC = -1.8369
-INDUS    => F = 1.0043, p = 0.3168, adjR2 = 0.7841, D_AIC = -0.9681
-ZN       => F = 4.5533, p = 0.0333, adjR2 = 0.7825, D_AIC = 2.6613
-CHAS     => F = 8.5584, p = 0.0036, adjR2 = 0.7808, D_AIC = 6.7263
-B        => F = 14.7979, p = 0.0001, adjR2 = 0.7780, D_AIC = 12.9946
-TAX      => F = 17.2838, p = 0.0000, adjR2 = 0.7770, D_AIC = 15.4705
-NOX      => F = 25.9206, p = 0.0000, adjR2 = 0.7732, D_AIC = 23.9797
-RAD      => F = 28.8640, p = 0.0000, adjR2 = 0.7719, D_AIC = 26.8471
-RM       => F = 29.4849, p = 0.0000, adjR2 = 0.7716, D_AIC = 27.4500
-DIS      => F = 37.8059, p = 0.0000, adjR2 = 0.7680, D_AIC = 35.4602
-PTRATIO  => F = 53.4154, p = 0.0000, adjR2 = 0.7611, D_AIC = 50.1529
-CRIM     => F = 60.9695, p = 0.0000, adjR2 = 0.7578, D_AIC = 57.1130
-LSTAT    => F = 204.5929, p = 0.0000, adjR2 = 0.6949, D_AIC = 173.9476
Removed AGE

Current model uses features ['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'DIS', 'RAD',
                             'TAX', 'PTRATIO', 'B', 'LSTAT']
-INDUS    => F = 1.0069, p = 0.3161, adjR2 = 0.7844, D_AIC = -0.9676
-ZN       => F = 4.4223, p = 0.0360, adjR2 = 0.7830, D_AIC = 2.5187
-CHAS     => F = 8.7145, p = 0.0033, adjR2 = 0.7811, D_AIC = 6.8661
-B        => F = 15.0808, p = 0.0001, adjR2 = 0.7783, D_AIC = 13.2464
-TAX      => F = 17.2296, p = 0.0000, adjR2 = 0.7774, D_AIC = 15.3819
-NOX      => F = 26.7816, p = 0.0000, adjR2 = 0.7732, D_AIC = 24.7672
-RAD      => F = 28.7545, p = 0.0000, adjR2 = 0.7723, D_AIC = 26.6841
-RM       => F = 31.7976, p = 0.0000, adjR2 = 0.7710, D_AIC = 29.6267
-DIS      => F = 42.9588, p = 0.0000, adjR2 = 0.7661, D_AIC = 40.2753
-PTRATIO  => F = 53.3780, p = 0.0000, adjR2 = 0.7616, D_AIC = 50.0178
-CRIM     => F = 61.0536, p = 0.0000, adjR2 = 0.7582, D_AIC = 57.0767
-LSTAT    => F = 227.5439, p = 0.0000, adjR2 = 0.6856, D_AIC = 190.0256

Final model uses 12 features ['ZN', 'INDUS', 'NOX', 'DIS', 'PTRATIO', 'CHAS', 'RM',
                              'TAX', 'B', 'LSTAT', 'CRIM', 'RAD']

```

We see that in the final pass, the  $F$ -score for the model without `INDUS` remains above the threshold  $F_{out}$ , even though the corresponding  $p$ -value is fairly high. Indeed, when using  $p$ -values to test hypotheses instead of the static  $F$ -score thresholds (note that the degrees of freedom change on each pass), all three methods produce the same 11 variable model (with a level of significance 0.05). With this, there seems to be no significant difference between the three procedures, apart from the number of iterations taken to reach the final model. The backward elimination model finishes much faster than the rest here, since the majority of variables here are significant; this may not always be the case if the available variables are not as suitable.

The code used to generate these models, as well as the output (using  $p$ -values, level of significance  $\alpha = 0.05$ ) has been attached.

From the procedure outputs, we see that all algorithms seem to follow the ‘path of least AIC’, i.e. among the available new models, the one with the greatest AIC drop is chosen. The procedure terminates when none of the new models offer an improvement via a drop in AIC.

Table 3: Parameter estimates using the models from Table 2.

Parameter	$k = 11$	$k = 12$	$k = 13$
$\beta_0$	4.0837	4.0952	4.1020
CRIM	-0.0103	-0.0103	-0.0103
ZN	0.0011	0.0011	0.0012
INDUS	—	0.0025	0.0025
CHAS	0.1051	0.1016	0.1009
NOX	-0.7217	-0.7623	-0.7784
RM	0.0907	0.0922	0.0908
AGE	—	—	0.0002
DIS	-0.0517	-0.0500	-0.0491
RAD	0.0134	0.0142	0.0143
TAX	-0.0006	-0.0006	-0.0006
PTRATIO	-0.0374	-0.0381	-0.0383
B	0.0004	0.0004	0.0004
LSTAT	-0.0286	-0.0288	-0.0290

## 2 Bootstrap methods

Given a table of  $n = 15$  schools, each with an LSAT and a GPA score, we can calculate the correlation between these two columns as  $\rho = 0.748$ . To estimate the standard error, we use the bootstrap procedure of drawing  $B$  random samples with replacement, each consisting of 15 rows from the data table. We calculate the correlation  $\rho_i$  between the LSAT and GPA for each sample  $1 \leq i \leq B$ , and thus obtain a bootstrap distribution  $\hat{f}_B$  for the correlation. Using this, we calculate an estimate of the standard error

$$[\hat{\text{se}}_B(\rho)]^2 = \frac{1}{B-1} \sum_{i=1}^B (\rho_i - \bar{\rho})^2.$$

The probability that  $\rho > 0.5$  can be estimated from this distribution as

$$\frac{1}{B} \sum_{i=1}^B I_{(0.5, 1]}(\rho_i),$$

i.e. the proportion of bootstrap estimates  $\rho_i$  that are above 0.5.

We perform these calculations for  $25 \leq B \leq 3000$  (reusing samples from the previous  $B$ ) and obtain the following variation of the estimate of  $\text{se}_B(\rho)$ .




Figure 1: Variation of  $\hat{se}_B(\rho)$  with  $B$

The final estimate of the standard error is 0.139 ( $B = 3000$ ).




Figure 2: Histogram of the 3000 bootstrap estimates  $\rho_i$ .

The estimate of  $P(\rho > 0.5)$  is 0.942.