

MA4106: Statistics II

Satvik Saha, 19MS154

October 15, 2022

1 Introduction

Here, we perform linear regression on the Boston house-price dataset; specifically, we regress the logarithm of **MEDV** (the median value of owner-occupied homes in units of \$1000) against 13 variables. The values Y_i , of the dependent variable, where i runs over the N row indices of data entries, have been assembled in the column vector \mathbf{Y} . Similarly, the values of X_{ij} of the independent variables, where i runs through the row indices and $j = 1, \dots, 13$ runs through the variable indices, have been assembled in the $N \times 14$ matrix \mathbf{X} . Note that the first column \mathbf{X}_0 has been filled with 1's. We now use the linear model

$$\mathbf{Y} = \mathbf{X}\beta + \epsilon,$$

where β is a column vector, and ϵ is a multivariate normal random variable with zero mean, covariance matrix $\sigma^2 \mathbb{I}_N$.




Figure 1: Correlation matrix between the different variables. Here, LOGMEDV denotes the (natural) logarithm of MEDV. Note the strong positive correlation between the variables RAD and TAX. This may lead to a problem of multicollinearity.

j	Variable	Description
1	CRIM	Per capita crime rate by town
2	ZN	Proportion of residential land zoned for lots over 25,000 sq.ft.
3	INDUS	Proportion of non-retail business acres per town
4	CHAS	Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
5	NOX	Nitric oxides concentration (parts per 10 million)
6	RM	Average number of rooms per dwelling
7	AGE	Proportion of owner-occupied units built prior to 1940
8	DIS	Weighted distances to five Boston employment centres
9	RAD	Index of accessibility to radial highways
10	TAX	Full-value property-tax rate per \$10,000
11	PTRATIO	Pupil-teacher ratio by town
12	B	$1000(Bk - 0.63)^2$ where Bk is the proportion of blacks by town
13	LSTAT	% lower status of the population
	MEDV	Median value of owner-occupied homes in \$1000's

Table 1: Variables and their descriptions.

Before performing regression, we normalize the independent variable values via

$$\mathbf{X}_j^* = \frac{\mathbf{X}_j - \bar{X}_j}{\sqrt{s_{jj}}}, \quad s_{jj} = \sum_{i=1}^N (\mathbf{X}_{ij} - \bar{X}_j)^2.$$

These are then assembled as the columns of the matrix \mathbf{X}_s^* . Then, our model becomes

$$\mathbf{Y} = \mathbf{X}_s^* \beta^* + \epsilon.$$

Here,

$$\beta_0 = \beta_0^* - \sum_{j=1}^{13} \frac{\beta_j^* \bar{X}_j}{\sqrt{s_{jj}}}, \quad \beta_j = \frac{\beta_j^*}{\sqrt{s_{jj}}}.$$

2 Ordinary Least Squares linear regression

Now, the Ordinary Least Squares estimate of β^* is given by

$$\hat{\beta}^* = (\mathbf{X}_s^{*\top} \mathbf{X}_s^*)^{-1} \mathbf{X}_s^{*\top} \mathbf{Y}.$$

Using this to predict the response given by the same N data entries gives the OLS prediction

$$\hat{\mathbf{Y}}_{\text{OLS}} = \mathbf{X}_s^* \hat{\beta}^* = \mathbf{X}_s^* (\mathbf{X}_s^{*\top} \mathbf{X}_s^*)^{-1} \mathbf{X}_s^{*\top} \mathbf{Y}.$$

The vector of residuals is given by

$$\mathbf{E} = \mathbf{Y} - \hat{\mathbf{Y}}_{\text{OLS}}.$$

The goodness of fit R^2 is given by

$$R^2 = 1 - \frac{\|\mathbf{E}\|^2}{\|\mathbf{Y} - \bar{Y}\|^2}.$$

Following are the results of OLS linear regression to the given dataset. This yields an R^2 score of 0.79.

j	Variable	$\hat{\beta}_j^*$	$\hat{\beta}_j$	Variance Inflation Factor
0		3.034513	4.102042	
1	CRIM	-1.985444	-0.010272	1.792192
2	ZN	0.614496	0.001172	2.298758
3	INDUS	0.380297	0.002467	3.991596
4	CHAS	0.575847	0.100888	1.073995
5	NOX	-2.026973	-0.778399	4.393720
6	RM	1.434196	0.090833	1.933744
7	AGE	0.133223	0.000211	3.100826
8	DIS	-2.322810	-0.049087	3.955945
9	RAD	2.791697	0.014267	7.484496
10	TAX	-2.370043	-0.000626	9.008554
11	PTRATIO	-1.861951	-0.038271	1.799084
12	B	0.848474	0.000414	1.348521
13	LSTAT	-4.659487	-0.029036	2.941491

Table 2: Parameters obtained via OLS linear regression on the Boston dataset.




Figure 2: Residuals obtained from the OLS linear regression.

Denoting \mathbf{X}^* to be the matrix obtained by deleting the first column (of 1's) from \mathbf{X}_s^* , note that $R_{xx} = \mathbf{X}^{*\top} \mathbf{X}^*$ is precisely the correlation matrix of the independent variables. Furthermore, the diagonal of R_{xx}^{-1} contains the Variance Inflation Factors of each of the variables, calculated in Table 2. The Variance Inflation Factor VIF_j of the j -th variable is so called since

$$\text{var}(\hat{\beta}_j^*) = \sigma^2 (\mathbf{X}^{*\top} \mathbf{X}^*)_{jj}^{-1} = \sigma^2 \text{VIF}_j.$$

High VIFs indicate that the estimates $\hat{\beta}_j^*$ for the corresponding variables are poor due to their high variance. Note the large VIFs for the variables **RAD** and **TAX**, which confirms our earlier suspicion that they introduce a problem of multicollinearity in our model.

3 Ridge regression

Here, our estimate of the parameters β^* is given by

$$\hat{\beta}^*(k) = (\mathbf{X}_s^{*\top} \mathbf{X}_s^* + k \mathbb{I}_{14})^{-1} \mathbf{X}_s^{*\top} \mathbf{Y}.$$

This is done to mitigate the problem of multicollinearity, which may cause the matrix $\mathbf{X}_s^{*\top} \mathbf{X}_s^*$ to be close to singular.

Let $\{\lambda_i\}$ be the eigenvalues of the matrix R_{xx} . Then, the value

$$\kappa = \sqrt{\frac{\lambda_{max}}{\lambda_{min}}}$$

is a measure of multicollinearity in the model. In our case, $\kappa \approx 9.82$.

3.1 Ridge-trace plot




Figure 3: A ridge-trace plot, showing the variation of the parameters β_j^* with k . It can be seen that each $\beta_j^* \rightarrow 0$ as $k \rightarrow \infty$. The vertical green line marks $k = 0.01$, and the vertical blue line marks $k = 0.008$. Note that **BIAS** denotes the ‘variable’ corresponding to the parameter β_0 .

The ridge trace plot shows that with increasing k , the parameters β_j^* shift from their OLS values and briefly plateau at another value, before decaying to zero. This effect is clearest in the curves of **TAX** and **ZN**, both around $k = 1$. Some parameters, like the curves for **INDUS**, **AGE**, **RAD**, **DIS**, even change sign. Note that the curves for **RAD** and **TAX** are the fastest shrinking ones with k .

By a crude visual estimate, we choose $k = 0.01$. This the point (to the nearest power of 10) where the some of the OLS estimates begin decaying towards zero.

Note that the OLS estimates of the parameters (Table 2) corresponding to `INDUS`, `AGE`, `RAD` are positive although they are negatively correlated with `LOGMEDV` as per Figure 1. In particular, the parameter corresponding to `RAD` has the third highest magnitude and is positive, despite the correlation between `RAD` and `LOGMEDV` being -0.48 . Similarly, the OLS estimate of the parameter corresponding to `DIS` has the fifth highest magnitude and is negative, despite the correlation between `DIS` and `LOGMEDV` being $+0.34$.

3.2 Cross validation

Another method of choosing a value of k is by minimizing the prediction error with respect to k . The Prediction Error (PE) has to be estimated via some means; here, we use Leave-One-Out Cross Validation (LOOCV). In this process, a row i of the dataset is set aside, and the model parameters are determined using the remaining rows. These are used to generate a prediction \hat{Y}_i for Y_i , using the variables in row i . This predicted value is compared against the true value of row i , yielding a prediction error. The above process is repeated for each row i . Here, we use the mean of squares of prediction errors,

$$\text{PE} \approx \text{LOOCV} = \frac{1}{N} \sum_{i=1}^N (Y_i - \hat{Y}_i)^2.$$




Figure 4: The variation of the LOOCV with k , which gives an estimate of the prediction error.

In our case, the LOOCV is minimized at $k = 0.008$. The LOOCV surpasses its OLS ($k = 0$) value beyond $k = 0.017$.

4 Conclusion

j	Variable	$\hat{\beta}_j^*(k)$			$\hat{\beta}_j(k)$		
		$k = 0$	$k = 0.008$	$k = 0.010$	$k = 0$	$k = 0.008$	$k = 0.010$
0		3.034513	3.034465	3.034453	4.102042	4.028287	4.010485
1	CRIM	-1.985444	-1.943415	-1.933262	-0.010272	-0.010054	-0.010002
2	ZN	0.614496	0.563474	0.551583	0.001172	0.001075	0.001052
3	INDUS	0.380297	0.278933	0.255974	0.002467	0.001809	0.001660
4	CHAS	0.575847	0.590940	0.594359	0.100888	0.103532	0.104131
5	NOX	-2.026973	-1.925292	-1.900518	-0.778399	-0.739352	-0.729838
6	RM	1.434196	1.490724	1.504286	0.090833	0.094413	0.095272
7	AGE	0.133223	0.090123	0.079812	0.000211	0.000142	0.000126
8	DIS	-2.322810	-2.243870	-2.223905	-0.049087	-0.047419	-0.046997
9	RAD	2.791697	2.495783	2.427773	0.014267	0.012755	0.012407
10	TAX	-2.370043	-2.098250	-2.037197	-0.000626	-0.000554	-0.000538
11	PTRATIO	-1.861951	-1.826501	-1.817834	-0.038271	-0.037543	-0.037365
12	B	0.848474	0.851608	0.852371	0.000414	0.000415	0.000415
13	LSTAT	-4.659487	-4.577014	-4.556287	-0.029036	-0.028522	-0.028392

Table 3: Parameter estimates for selected values of k .

The R^2 values corresponding to the ridge regressions with $k = 0, 0.008, 0.010$ have been calculated as 0.7896, 0.7894, 0.793 respectively. Figure 4 shows that the estimated prediction error for $k = 0.010$ is slightly greater than that for $k = 0.008$. Thus, we choose the ridge regression parameter 0.008, as it minimizes the estimated prediction error. This choice does not contradict the ridge-trace plot either.

Also note that we have calculated $\kappa < 10$, which does not seem to indicate much of a problem with multicollinearity. The main variables under suspicion here are **RAD** and **TAX** which are highly correlated; simply removing one of the two may yield a better performing model than one obtained via ridge regression. Removing some of the variables which change sign in the ridge-trace plot (Figure 3) may also yield some improvement.

With this in mind, we also prefer the ridge estimates to the OLS ones; the ridge estimates give better (estimated) prediction errors, with lower variance inflations.