

MA4206: Linear Models

Satvik Saha, 19MS154

March 23, 2023

1 Introduction

There are 33 countries with data on unemployment rates, GDP per capita, and total investment as a percentage of GDP, denoted UNMP (y), GDP (x_1), INV (x_2) respectively.

These are fitted against the linear model

$$y_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \epsilon_i. \quad (\star)$$

The fitted parameters are as follows, with an adjusted R^2 value of 0.477.

Table 1: Fitted parameters from the linear model (\star).

	Variable	Value	Standard error
Intercept	β_0	2.639×10^1	0.347×10^1
GDP	β_1	-7.926×10^{-5}	3.204×10^{-5}
INV	β_2	-7.051×10^{-1}	1.598×10^{-1}

We further plot fitted values of UNMP vs observed values of UNMP.




Figure 1: UNMP values, fitted vs actual. The blue line marks $x = y$.

From the above plot, we see that the model performs especially badly on the two points on the extreme right, which represent the countries Greece and Spain. Elsewhere, the model seems to explain the general trend of observed values, albeit poorly. Thus, the goodness of fit is poor.

Plotting the residuals vs fitted UNMP values further highlights that Greece and Spain are outliers (with the highest residual values). Although $\hat{\epsilon}$ and \hat{y} are uncorrelated, we calculate a Spearman correlation of -0.24.




Figure 2: Residuals vs fitted values of UNMP.

2 Prediction intervals

The 95% prediction intervals for the model have been calculated and displayed below.




Figure 3: Prediction intervals for UNMP. The blue points mark the fitted values, while the red points mark the observed values.

3 Leave One Out Cross Validation

We perform Leave One Out Cross Validation (LOOCV) by fitting the model against the data with the i -th entry removed, for each of the 33 rows indexed by i . The tuples of fitted parameters $(\beta_0^{(i)}, \beta_1^{(i)}, \beta_2^{(i)})$ have been plotted in space, with the points labelled by the i -th country.

Table 2: Fitted parameters from the linear model (\star) with the i -th entry removed.

i	Country	$\beta_0^{(i)}$	$\beta_1^{(i)}$	$\beta_2^{(i)}$
1	Australia	26.826	-0.0000805	-0.727
2	Austria	26.355	-0.0000792	-0.703
3	Belgium	26.457	-0.0000791	-0.711
4	Canada	26.441	-0.0000793	-0.708
5	Cyprus	26.319	-0.0000791	-0.702
6	CzechRepublic	26.257	-0.0000810	-0.694
7	Denmark	26.558	-0.0000767	-0.715
8	Estonia	26.282	-0.0000827	-0.691
9	Finland	26.338	-0.0000795	-0.704
10	France	26.394	-0.0000785	-0.711
11	Germany	26.985	-0.0000779	-0.728
12	Greece	21.339	-0.0000712	-0.498
13	Iceland	26.955	-0.0000739	-0.734
14	Ireland	26.431	-0.0000849	-0.701
15	Israel	26.809	-0.0000802	-0.717
16	Italy	26.366	-0.0000792	-0.704
17	Japan	26.184	-0.0000815	-0.686
18	Korea	26.328	-0.0000797	-0.701
19	Latvia	26.382	-0.0000789	-0.706
20	Lithuania	26.688	-0.0000837	-0.708
21	Luxembourg	26.345	-0.0000965	-0.676
22	Netherlands	26.664	-0.0000780	-0.717
23	NewZealand	26.355	-0.0000799	-0.699
24	Norway	27.688	-0.0000892	-0.753
25	Portugal	26.902	-0.0000813	-0.723
26	Singapore	26.093	-0.0000783	-0.691
27	SlovakRepublic	26.398	-0.0000732	-0.721
28	Slovenia	26.582	-0.0000817	-0.707
29	Spain	25.068	-0.0000666	-0.685
30	Sweden	26.590	-0.0000806	-0.715
31	Switzerland	26.330	-0.0000772	-0.705
32	UnitedKingdom	27.505	-0.0000755	-0.756
33	UnitedStates	26.514	-0.0000759	-0.714

If the fitted parameters for some i deviate significantly from the rest, we hypothesize that the corresponding country heavily influences our model. Indeed, the countries Greece and Spain deviate the most.

This degree of outlyingness can be loosely quantified by the Mahalanobis depth, also indicated in a figure below.

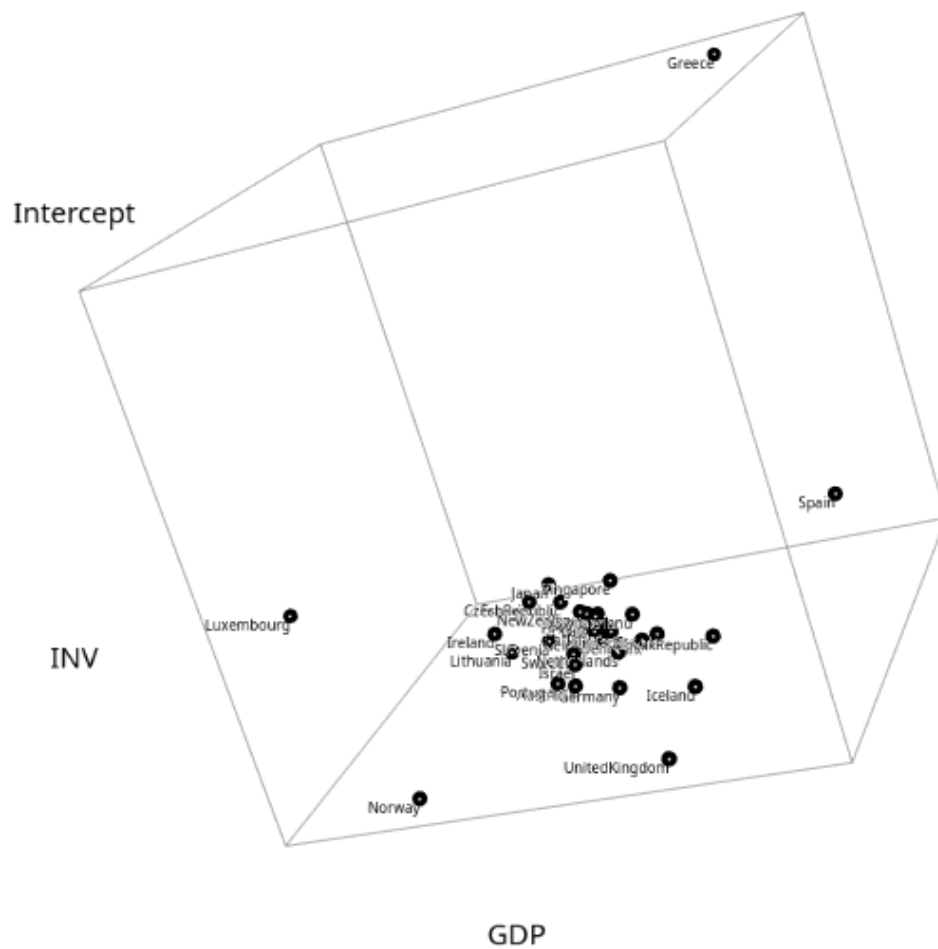


Figure 4: Fitted parameter values obtained from the linear model upon removing the entry with the labelled country.




Figure 5: Mahalanobis depths of the data points in Figure 4.

4 Removing outliers

We removed the rows corresponding to the countries Greece and Spain, and perform the same procedures as before. Our new parameter estimates are given below, with an adjusted R^2 of 0.415.

Table 3: Fitted parameters from the linear model (\star).

	Variable	Value	Standard error
Intercept	β_0^*	1.923×10^1	0.272×10^1
GDP	β_1^*	-5.655×10^{-5}	2.124×10^{-5}
INV	β_2^*	-4.481×10^{-1}	-3.717×10^{-1}

The fitted vs observed plot now exhibits a much more reasonable trend.




Figure 6: UNMP values, newly fitted vs actual. The blue line marks $x = y$.

The residuals now exhibit a Spearman correlation of -0.04 with the newly fitted UNMP values.




Figure 7: New residuals vs newly fitted values of UNMP.




Figure 8: Prediction intervals for UNMP. The blue points mark the fitted values, while the red points mark the observed values.




Figure 9: New fitted parameter values obtained from the linear model upon removing the entry with the labelled country.

5 Discussion

We have justified the removal of the countries Greece and Spain on the basis of their position in the fitted vs observed and residual vs fitted plots, as well as their outlyingness in the LOOCV fitted parameter scatter plot. With this, the remaining countries fit reasonably well in the model; the next candidate countries for removal would be Cyprus and Luxembourg.

These two countries exhibit an unemployment rate that is significantly higher than predicted. Thus, it indicates that their economics scenarios are very different from those of the remaining countries, which can be somewhat accurately modelled using only two variables.




Figure 10: Sorted values of GDP, INV, and UNMP.

We note that Greece has the lowest total investment rate by far, followed by Cyprus and Portugal, which may help explain Greece's status as an outlier. Additionally, Luxembourg has the highest GDP by far, followed by Switzerland and Norway.